

Aakash Shahani

aakashshahani21@gmail.com | linkedin.com/in/aakash-shahani | github.com/aakashshahani | [Portfolio](#)

EDUCATION

University of South Florida (USF)

Bachelor of Science in Computer Science

Graduation: December 2025

GPA: 3.5/4.0

TECHNICAL SKILLS

Languages: Python, Go, C/C++, Java, TypeScript, JavaScript, SQL, Bash

AI / ML: PyTorch, transformers, Hugging Face, scikit-learn, XGBoost, RAG, Embeddings, LangChain, Prompt Engineering, LLM APIs (Anthropic/OpenAI), SHAP

Data & Cloud: Airflow, dbt, Iceberg, Spark, Kafka, star schema & SCD2, Great Expectations, AWS

(S3/Glue/Athena), Terraform, Docker, CI/CD (GitHub Actions), pandas, NumPy, A/B Testing, Prometheus/Grafana

Backend, Databases & Tools: FastAPI, REST, GraphQL, Node.js, Next.js, React, PostgreSQL, MySQL, MongoDB, Redis, pgvector/FAISS, Metabase, Git, Linux, pytest

EXPERIENCE

AI/ML Engineer

Jan 2026 – Present

CSSAI Lab, University of South Florida

Tampa, FL

- Built a reproducible LLM evaluation pipeline (temperature-0 batch inference) benchmarking **3** frontier LLMs (GPT-4.1, Gemini 2.5 Flash, Claude Sonnet 4.6) against **6** trained baselines on **2** detection tasks.
- Drove the detector to **0.79** macro-F1 / **0.81** accuracy, significantly beating every trained baseline (XGBoost, fine-tuned DistilBERT) and approaching human-expert agreement, through precision-first prompt engineering, grouped cross-validation, and bootstrapped significance testing.

Machine Learning Research Intern

Aug 2024 – Apr 2025

University of South Florida

Tampa, FL

- Trained an LSTM binding-site classifier (focal loss, Adam, dropout, LR scheduling) to **0.65** AUPRC / **0.61** MCC on a held-out test set, benchmarked against published CLAPE-SMB/SCRIBER with k-fold cross-validation.
- Exposed **46%** train-test homology leakage in the published binding-site benchmark via MMseqs2 **30%**-identity clustering, re-scoring with cluster-grouped k-fold CV for an honest **0.57** AUPRC.

PUBLICATIONS

R. A. Baten, M. A. Urmi, **A. Shahani**, “Strategic Adaptability Under Changing Contexts: Behavioral Reorganization in Negotiation,” *Computers in Human Behavior* (Elsevier), under review, 2026.

PROJECTS

ScholarLens: Agentic Research Platform | FastAPI, Next.js, pgvector, Claude API

[GitHub](#) | [Live Demo](#)

- Built an agentic RAG platform that flags cross-paper contradictions via a two-stage BM25 + dense pre-filter and an LLM judge at macro-F1 **0.788** / κ **0.683**, cutting $\sim$ **4,500** LLM calls by over **90%**.
- Ran **6** analyses per paper concurrently for a **5 $\times$** speedup and served pgvector (HNSW) search behind a cross-encoder reranker at nDCG@5 **0.98**; deployed live with multi-user isolation.

TradePulse: Real-Time Streaming Pipeline | Kafka, Spark Streaming, Docker

[GitHub](#)

- Streamed live Coinbase trades through a **6**-partition Kafka topic into 1-minute OHLC candles in Spark Structured Streaming, with exactly-once output via watermarks and idempotent writes (crash test: **0** duplicates).
- Sustained sub-second **600–800ms** micro-batches under $\sim$ **1,500** rows/sec bursts across **9** Dockerized services with Prometheus/Grafana observability; validated by **20** tests in CI.

NewsRec: Cold-Start News Recommender | PyTorch, FAISS, FastAPI, Docker

[GitHub](#)

- Trained a content-based two-tower retriever on the MIND dataset that beat popularity/SVD/ALS baselines by **+0.124** AUC ($\sim$ **23%**) across **73,152** impressions by embedding article text to solve cold-start.
- Served candidate generation at **2.8ms** p99 over **3,000** users via a FAISS HNSW index so the FastAPI path never loads a model at request time; Docker + **17** tests in CI.

UrbanFlow: Data Lakehouse Platform | Airflow, dbt, Apache Iceberg, AWS, Terraform

[GitHub](#)

- Built an end-to-end data lakehouse over **3M** NYC taxi trips: Airflow into Apache Iceberg on S3, dbt-modeled into a star schema (fact + SCD Type 2 dims), deployed on AWS with Terraform and queried on Athena/Trino.
- Cut Athena bytes scanned **97%** (**4.81** $\rightarrow$ **0.13MB**) via day-partitioning; gated it with Great Expectations and **17** dbt tests that block bad data before the gold layer; idempotent Iceberg MERGE held **2.87M** rows at **0** duplicates.